

## Homework 2 – Deep Neural Networks (CS525 191D, Whitehill, Spring 2018)

You may complete this homework assignment either individually or in teams up to 2 people.

1. **XOR problem** [10 points]: Show (by deriving the gradient, setting to 0, and solving mathematically, not in Python) that the values for  $\mathbf{w} = (w_1, w_2)$  and  $b$  that minimize the function  $J(\mathbf{w}, b)$  in Equation 6.1 (in the *Deep Learning* textbook) are:  $w_1 = 0$ ,  $w_2 = 0$ , and  $b = 0.5$ . Put your solution in a PDF file called `homework2.WPIUSERNAME1.pdf` (or `homework2.WPIUSERNAME1.WPIUSERNAME2.pdf` for teams).
2. **Smile detector**: Train a simple “smile detector” that analyzes a  $(24 \times 24 = 576)$ -pixel grayscale face image and outputs a real number  $\hat{y}$  representing whether or not the image is smiling ( $\hat{y}$  close to 1 means “smile”;  $\hat{y}$  close to 0 means “non-smile”). Your detector should be implemented as a neural network  $f_{\mathbf{w}} : \mathbb{R}^{576} \rightarrow \mathbb{R}$  consisting of just an input layer and an output layer, with no layers in between. (This network is therefore not very “deep”, but you have to start somewhere.) **Note 1**: you must complete this problem using only linear algebraic operations in `numpy` – you may **not** use any off-the-shelf linear regression or neural network training software, as that would defeat the purpose. **Note 2**: If you have OpenCV 2.4.13 or higher, you can run a real-time demo (uncomment the corresponding lines in `homework2.template.py`) of the smile detector you train.
  - (a) **Method 1 – set gradient to 0 and solve** [12 points]: Compute the parameters  $\mathbf{w} = (w_1, \dots, w_{576})$  representing the “weights”/parameters of the neural network by deriving the expression for the gradient of the cost function w.r.t.  $\mathbf{w}$ , setting it to 0, and then solving. The cost function is

$$J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2$$

where  $\hat{y} = f_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^\top \mathbf{x}$  and  $m$  is the number of examples in the training set  $\mathcal{D}_{\text{tr}} = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(m)}, y^{(m)})\}$ , each  $\mathbf{x}^{(i)} \in \mathbb{R}^{576}$  and each  $y^{(i)} \in \{0, 1\}$ . Note that this “one-shot” method of optimizing the neural network parameters only works in **very particular cases** (such as linear regression, as we have here). After optimizing  $\mathbf{w}$  only on the **training set**, compute and report the cost  $J$  on the training set  $\mathcal{D}_{\text{tr}}$  and (separately) on the testing set  $\mathcal{D}_{\text{te}}$ .

- (b) **Method 2 – gradient descent** [12 points]: Pick a random starting value for  $\mathbf{w} \in \mathbb{R}^{576}$  and a small learning rate ( $\epsilon \ll 1$ ). Then, using the expression for the gradient of the cost function, iteratively update  $\mathbf{w}$  to reduce the cost  $J(\mathbf{w})$ . Stop when the difference between  $J$  over successive training rounds is below some “tolerance” (e.g.,  $\delta = 0.001$ ). After optimizing  $\mathbf{w}$  only on the **training set**, compute and report the cost  $J$  on the training set  $\mathcal{D}_{\text{tr}}$  and (separately) on the testing set  $\mathcal{D}_{\text{te}}$ . Both of these values should be very close to what you computed using Method 1. Note that this method of optimizing neural network parameters is **much more general** than Method 1 above, at the expense of requiring some additional optimization hyperparameters (e.g., learning rate, tolerance).
- (c) **Method 3 – gradient descent with regularization** [6 points]: Same as (b) above, but change the cost function to include a penalty for  $\|\mathbf{w}\|^2$  growing too large:

$$\tilde{J}(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2 + \frac{\alpha}{2} \mathbf{w}^\top \mathbf{w}$$

where  $\alpha \in \mathbb{R}^+$ . Set  $\alpha = 1000$  and then optimize  $\tilde{J}$  w.r.t.  $\mathbf{w}$ . After optimizing  $\mathbf{w}$  only on the **training set** (using  $\tilde{J}$ ), compute and report the *unregularized* cost  $J$  on the training set  $\mathcal{D}_{\text{tr}}$  and (separately) the testing set  $\mathcal{D}_{\text{te}}$ . The training cost should be higher (i.e., worse), but the testing cost should be lower (i.e., better). How does the value of  $\|\mathbf{w}\|^2$  using Method 3 compare to its value using Method 2?

Put your solution in a Python file called `homework2.WPIUSERNAME1.py`  
(or `homework2.WPIUSERNAME1.WPIUSERNAME2.py` for teams).