

## Homework 6 – Deep Neural Networks (CS525 191D, Whitehill, Spring 2018)

You may complete this homework assignment either individually or in teams up to 2 people.

In this assignment you will train a recurrent neural network to solve a simple simulation task ( $n$ -back problem). The RNN you create and train should adhere to the following specification:

$$\begin{aligned} J_t(\mathbf{U}, \mathbf{V}, \mathbf{w}) &= \frac{1}{2}(\hat{y}_t - y_t)^2 \\ \hat{y}_t &= \mathbf{h}_t^\top \mathbf{w} \\ \mathbf{h}_0 &= \mathbf{0} \\ \mathbf{h}_t &= \tanh(\mathbf{z}_t) \\ \mathbf{z}_t &= \begin{bmatrix} \mathbf{U} & \mathbf{V} \end{bmatrix} \begin{bmatrix} \mathbf{h}_{t-1} \\ \mathbf{x}_t \end{bmatrix} \\ &= \begin{bmatrix} \mathbf{U}\mathbf{h}_{t-1} + \mathbf{V}\mathbf{x}_t \end{bmatrix} \end{aligned}$$

For simplicity, each  $y_t \in \mathbb{R}$  (i.e., a scalar). We will use the sum-of-squared-errors cost function.

Your job is to (1) derive the gradient updates for the network; and (2) implement SGD for this procedure to solve the  $n$ -back problem (for just  $n = 2$ ). In particular, for a sequence containing  $T$  terms (with inputs  $\mathbf{x}_1, \dots, \mathbf{x}_T$  and corresponding outputs  $y_1, \dots, y_T$ ), you should compute the gradient, w.r.t. each of  $\mathbf{U}, \mathbf{V}, \mathbf{w}$ , of the *total* cost  $J(\mathbf{U}, \mathbf{V}, \mathbf{w}) = \sum_{t=1}^T J_t(\mathbf{U}, \mathbf{V}, \mathbf{w})$ .

1. **Gradient derivation (30 points):** Derive **step-by-step** using the chain rule of matrix/vector calculus the three gradient terms; apply whatever analytical simplifications you can. Your math derivation should yield a **simple** algorithm for how to update each of these parameters in Python code.

$$\begin{aligned} \frac{\partial J}{\partial \text{vec}[\mathbf{U}]} &= \dots \\ \frac{\partial J}{\partial \text{vec}[\mathbf{V}]} &= \dots \\ \frac{\partial J}{\partial \mathbf{w}} &= \dots \end{aligned}$$

2. **Application to  $n$ -back problem (30 points):** Implement your math derivations above in Python code to conduct SGD on an RNN containing 6 hidden units, 1 input unit, and 1 output unit. **Include a screenshot** showing how the loss decreases during training. There is no validation or test set in this assignment – just focus on driving the loss toward 0. Depending on the initial seed of the data generator, you should be able to obtain a loss value  $< 0.05$ , e.g.:

11.4657453617  
11.2834015915  
11.079596572  
10.849944084  
10.5912803802  
10.3018414418  
9.98134933669  
9.63088887735  
9.25253518532  
8.84889117335  
8.42302162225

```
7.97962392769
7.52826431536
...
0.0530270146118
0.0499058526839
```

Note that, even with correctly implemented gradient expressions, the cost function may not decrease monotonically (depending on the learning rate). For this reason, it is strongly recommended that you use `check_grad` or `approx_fprime` to help you debug your gradient descent procedure.

In addition to your Python code (`homework6_WPIUSERNAME1.py` or `homework6_WPIUSERNAME1_WPIUSERNAME2.py` for teams), create a PDF file (`homework6_WPIUSERNAME1.pdf` or `homework6_WPIUSERNAME1_WPIUSERNAME2.pdf` for teams) containing the screenshots described above.